Statistical inference and random sampling

FOUNDATIONS OF INFERENCE IN PYTHON

Paul Savala
Assistant Professor of Mathematics

datacamp

Descriptive statistics

- Sample statistics meant to summarize the data
- Descriptive statistics summarize our sample

Date	SP500 Close	Daily Change
2017-08-07	2480.91	6.14
2017-08-08	2474.92	-5.99
2017-08-09	2474.02	-0.90
2017-08-10	2438.21	-35.81

Average daily change: -$9.14

Inference

- Infer something about our population
- **Descriptive statistics:** Describe data
- **Inference:** Make conclusions and decisions

Date	SP500 Close	Daily Change
2017-08-07	2480.91	6.14
2017-08-08	2474.92	-5.99
2017-08-09	2474.02	-0.90
2017-08-10	2438.21	-35.81

Average daily swing for *any* days ~$9.14

Statistical inference process

STATISTIC

Compute a **statistic** using our sample

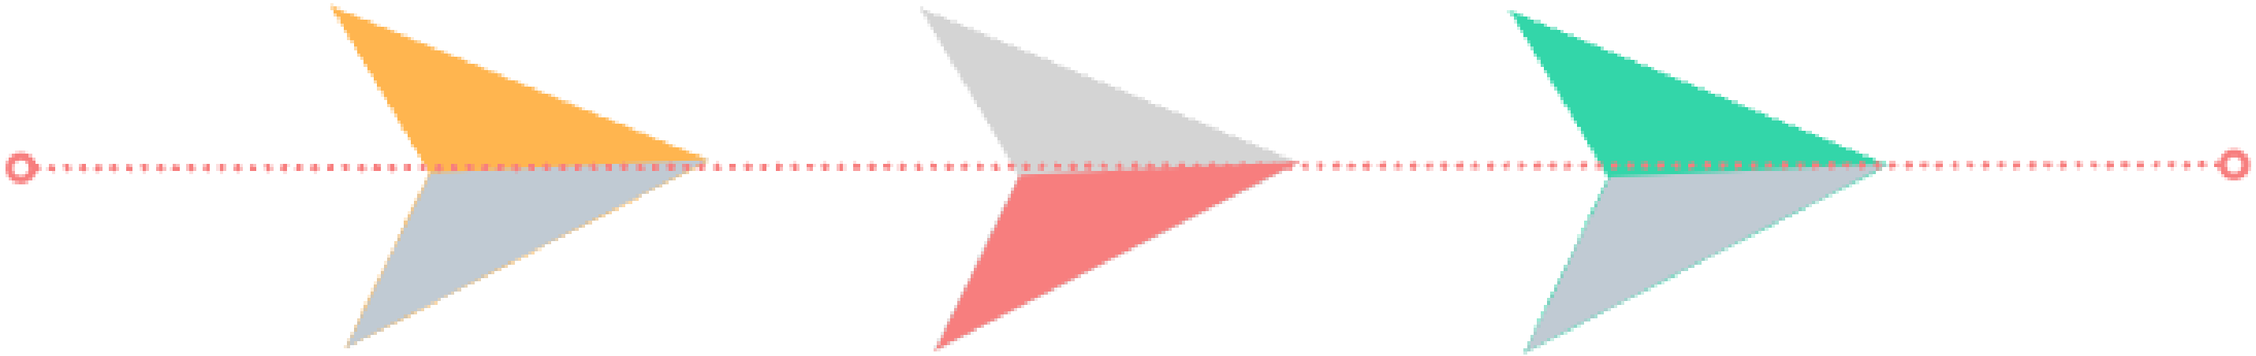

SAMPLE

Start by collecting a **sample** of data from the population

INFERENCE

Make inference about corresponding

Point estimates

- Given by a single value
- "Best guess" at an unknown population statistic

Point estimate: 1158.95 BTC daily swing

	Date	High_BTC	Low_BTC
0	2017-08-07	3397.679932	3180.889893
1	2017-08-08	3484.850098	3345.830078
2	2017-08-09	3422.760010	3247.669922
3	2017-08-10	3453.449951	3319.469971
4	2017-08-11	3679.719971	3372.120117

```
btc_high = btc_sp_df['High_BTC']
btc_low = btc_sp_df['Low_BTC']

np.mean(btc_high - btc_low)
```

```
1158.95
```

Sampling

Point estimates depend on the sample

```
btc_sp_first100 = btc_sp_df.iloc[:100]
np.mean(btc_sp_first100['High_BTC'] - btc_sp_first100['Low_BTC'])
```

```
659.60
```

```
initial_row = np.random.choice(btc_sp_df.shape[0]-100)
btc_sp_random_100 = btc_sp_df.iloc[initial_row:initial_row+100]
np.mean(btc_sp_first100['High_BTC'] - btc_sp_first100['Low_BTC'])
```

```
943.83
```

Let's practice!

FOUNDATIONS OF INFERENCE IN PYTHON

Sampling and bias

FOUNDATIONS OF INFERENCE IN PYTHON

Paul Savala
Assistant Professor of Mathematics

datacamp

Bias

- **Biased sample:** A group occurs more/less often in sample than in population

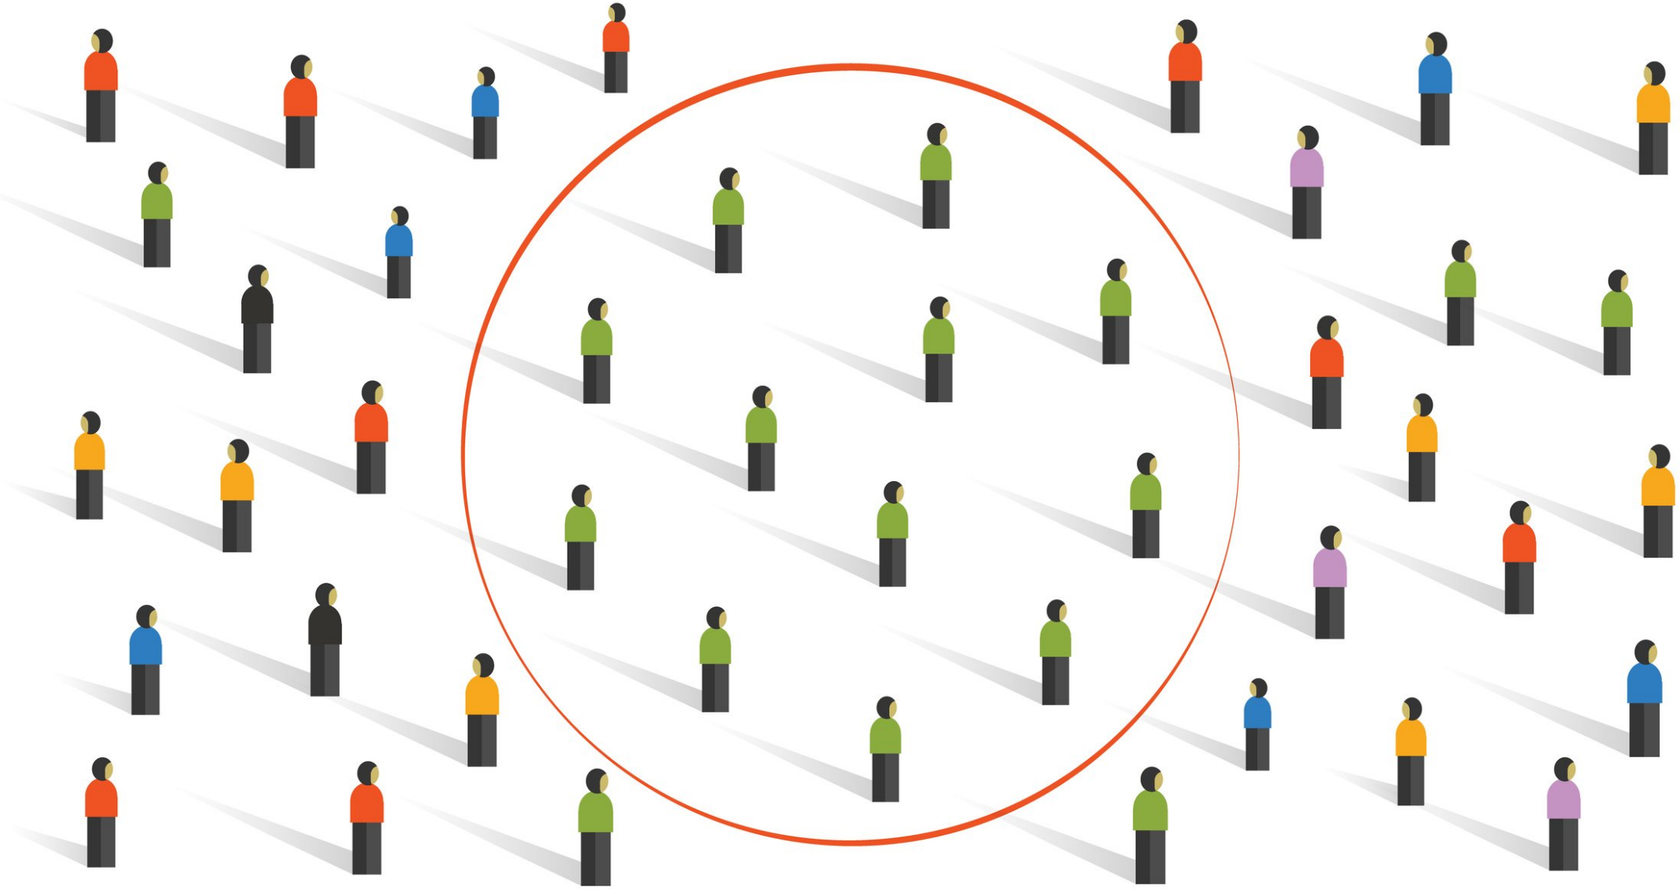

Biased samples

```
all_salaries = [75000, 82000, ...]
friends_salaries = [93000, 87000, 103000, 101000]


np.mean(friends_salaries)
```

```
96000
```

Sampling distribution

```
sampling_distribution = []

for i in range(100):
    random_sample = np.random.choice(salaries, size=10)
    sample_mean = np.mean(random_sample)
    sampling_distribution.append(sample_mean)

plt.hist(sampling_distribution)
plt.xlabel('Mean salary')
plt.ylabel('Percent of samples')
plt.title('Sampling distribution of mean salaries')
plt.show()
```

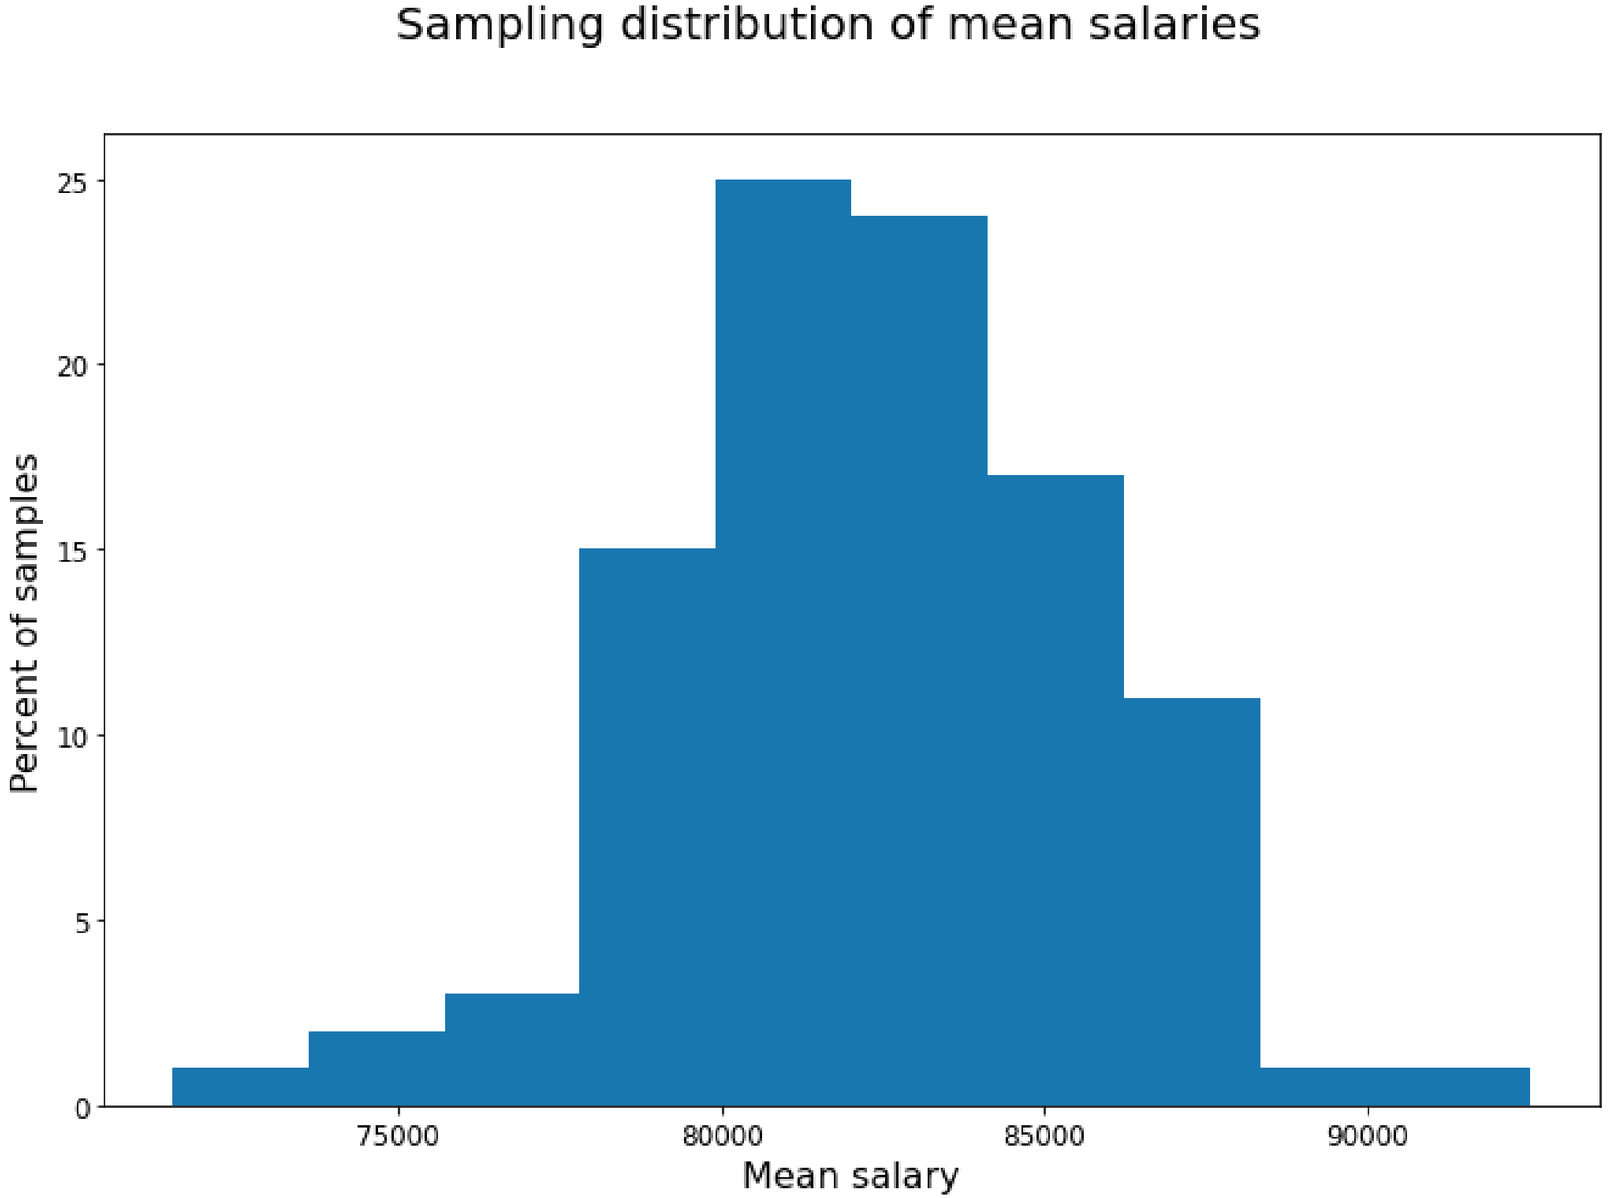
Sampling distribution of mean salaries
Percent of samples
0
5
10
15
20
25
75000
80000
85000
90000
Mean salary

Depends on the sample

- Samples affect point estimates
- Point estimates affect inference
- Samples affect p-value calculations

Doesn't depend on the sample

- Population statistic
  - Is unaffected by sample chosen
- Conclusion from test
  - Given a p-value, conclusion is unaffected by sample chosen

Let's practice!

FOUNDATIONS OF INFERENCE IN PYTHON

Confidence intervals and sampling

FOUNDATIONS OF INFERENCE IN PYTHON

Paul Savala
Assistant Professor of Mathematics

datacamp

What is a confidence interval?

- Uses samples to generate range of values
- Range of values estimate the population statistic

Example:

- Sample of 100 employees
- Mean salary of $80,000
- Standard deviation of $10,000

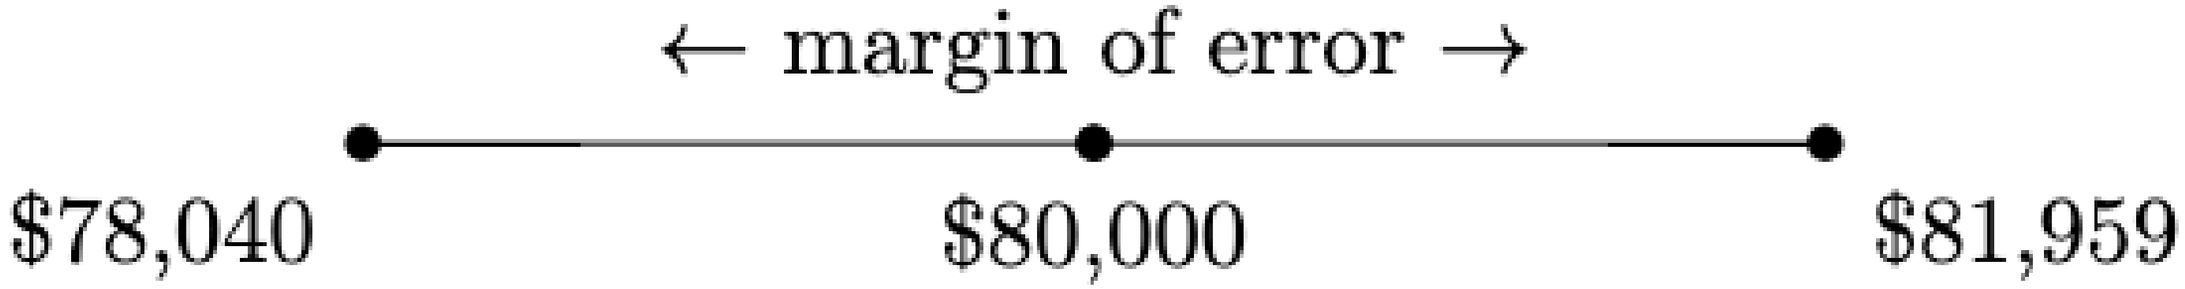

Calculating a confidence interval

```
from scipy import stats
import numpy as np

ci = stats.norm.interval(loc=80000,                    # Mean
                         scale=10000/np.sqrt(100),     # Standard error
                         alpha=0.95)                   # Confidence level
print(ci)
```

```
(78040.04, 81959.96)
```

Valid inference requires a normal sampling distribution

Central Limit Theorem

- Average many independent samples
- Sampling distribution is approximately normal

```
population = [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
sample_means = []

for i in range(1000):
  sample_5 = np.random.choice(population, size=5)
  sample_means.append(sample_5.mean())
```

```
plt.hist(sample_means)
```

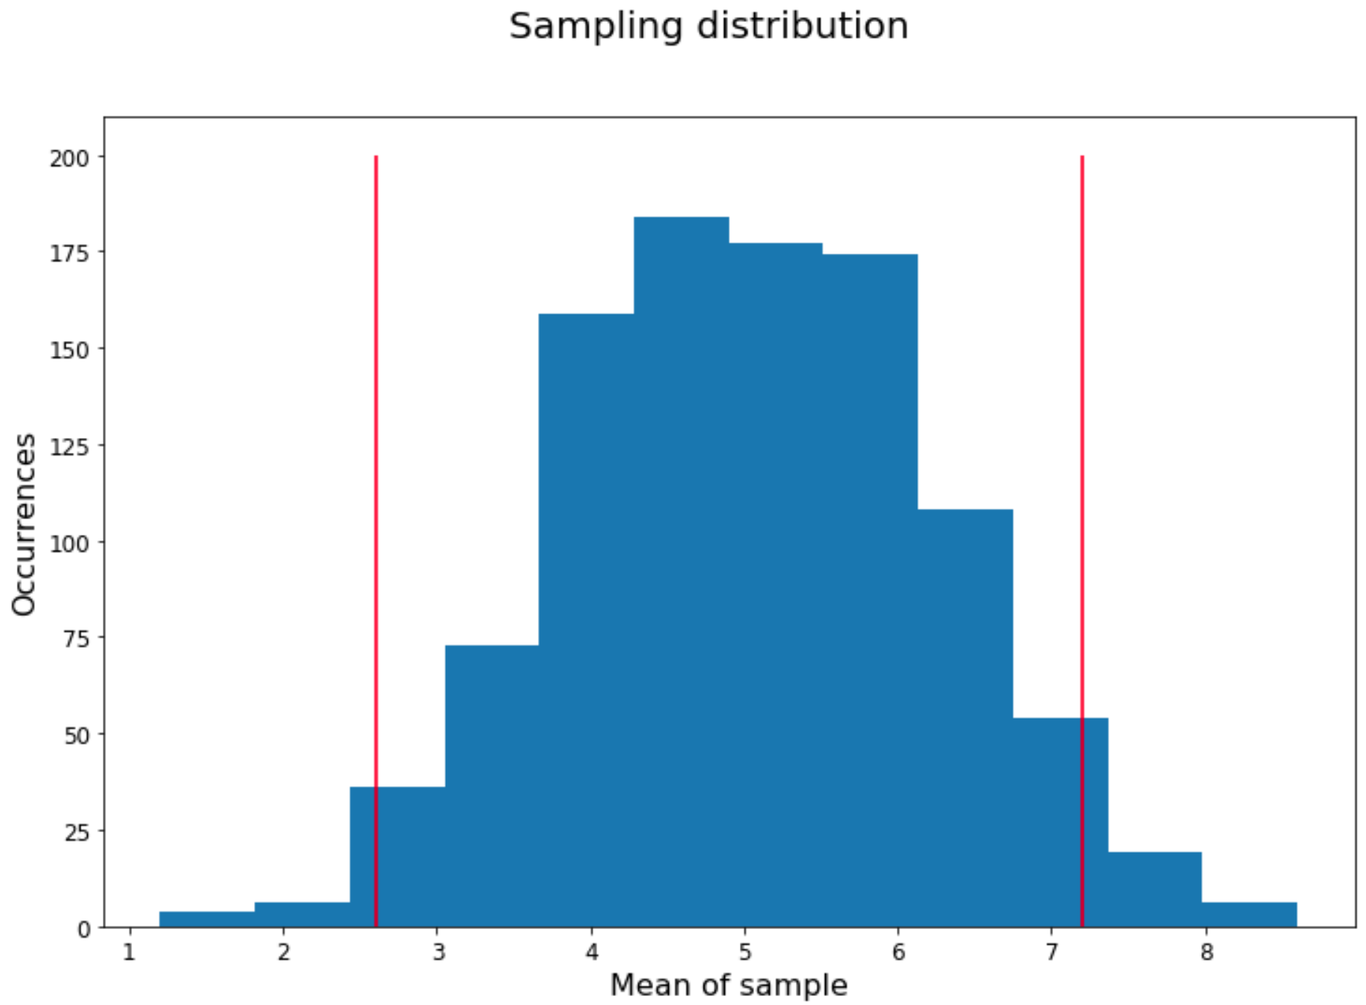

What a confidence interval tells us

(and what it doesn't tell us)

- Population statistic is or is not in confidence interval
- Repeated samples -> 95% of confidence intervals contain population statistic

Let's practice!

FOUNDATIONS OF INFERENCE IN PYTHON