


Introduction to preprocessing

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman

Curriculum Manager, DataCamp

What is data preprocessing?

- After exploratory data analysis and data cleaning
- Preparing data for modeling
- **Example:** transforming categorical features into numerical features (dummy variables)

Why preprocess?

- Transform dataset so it's suitable for modeling
- Improve model performance
- Generate more reliable results



Recap: exploring data with pandas

```
import pandas as pd  
hiking = pd.read_json("hiking.json")  
print(hiking.head())
```

```
Prop_ID          Name  ...  lat  lon  
0    B057  Salt Marsh Nature Trail  ...   NaN  NaN  
1    B073        Lullwater  ...   NaN  NaN  
2    B073      Midwood  ...   NaN  NaN  
3    B073  Peninsula  ...   NaN  NaN  
4    B073  Waterfall  ...   NaN  NaN
```

Recap: exploring data with pandas

```
print(hiking.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 33 entries, 0 to 32
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
 --   --  
 0   Prop_ID          33 non-null    object  
 1   Name              33 non-null    object  
 2   Location          33 non-null    object  
 3   Park_Name         33 non-null    object  
 4   Length             29 non-null    object  
 5   Difficulty         27 non-null    object  
 6   Other_Details      31 non-null    object  
 7   Accessible         33 non-null    object  
 8   Limited_Access     33 non-null    object  
 9   lat                0 non-null    float64 
 10  lon                0 non-null    float64 
dtypes: float64(2), object(9)
memory usage: 3.0+ KB
```

Recap: exploring data with pandas

```
print(wine.describe())
```

	Type	Alcohol	...	Alcalinity of ash
count	178.000000	178.000000	...	178.000000
mean	1.938202	13.000618	...	19.494944
std	0.775035	0.811827	...	3.339564
min	1.000000	11.030000	...	10.600000
25%	1.000000	12.362500	...	17.200000
50%	2.000000	13.050000	...	19.500000
75%	3.000000	13.677500	...	21.500000
max	3.000000	14.830000	...	30.000000

Removing missing data

```
print(df)
```

```
      A      B      C  
0  1.0    NaN  2.0  
1  4.0  7.0  3.0  
2  7.0    NaN    NaN  
3  NaN  7.0    NaN  
4  5.0  9.0  7.0
```

```
print(df.dropna())
```

```
      A      B      C  
1  4.0  7.0  3.0  
4  5.0  9.0  7.0
```

Removing missing data

```
print(df)
```

```
      A      B      C  
0  1.0    NaN  2.0  
1  4.0  7.0  3.0  
2  7.0    NaN    NaN  
3  NaN  7.0    NaN  
4  5.0  9.0  7.0
```

```
print(df.drop([1, 2, 3]))
```

```
      A      B      C  
0  1.0    NaN  2.0  
4  5.0  9.0  7.0
```

Removing missing data

```
print(df)
```

```
      A      B      C  
0  1.0    NaN  2.0  
1  4.0  7.0  3.0  
2  7.0    NaN    NaN  
3  NaN  7.0    NaN  
4  5.0  9.0  7.0
```

```
print(df.drop("A", axis=1))
```

```
      B      C  
0    NaN  2.0  
1  7.0  3.0  
2    NaN    NaN  
3  7.0    NaN  
4  9.0  7.0
```

Removing missing data

```
print(df)
```

```
A      B      C  
0    1.0    NaN    2.0  
1    4.0    7.0    3.0  
2    7.0    NaN    NaN  
3    NaN    7.0    NaN  
4    5.0    9.0    7.0
```

```
print(df.isna().sum())
```

```
A      1  
B      2  
C      2  
dtype: int64
```

```
print(df.dropna(subset=["B"]))
```

```
A      B      C  
1    4.0    7.0    3.0  
3    NaN    7.0    NaN  
4    5.0    9.0    7.0
```

Removing missing data

```
print(df)
```

```
      A      B      C  
0  1.0    NaN  2.0  
1  4.0  7.0  3.0  
2  7.0    NaN    NaN  
3  NaN  7.0    NaN  
4  5.0  9.0  7.0
```

```
print(df.dropna(thresh=2))
```

```
      A      B      C  
0  1.0    NaN  2.0  
1  4.0  7.0  3.0  
4  5.0  9.0  7.0
```

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Working With Data Types

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman

Curriculum Manager, DataCamp

Why are types important?

```
print(volunteer.info())
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 665 entries, 0 to 664  
Data columns (total 35 columns):  
 #   Column           Non-Null Count  Dtype     
 --  --  
 0   opportunity_id  665 non-null    int64    
 1   content_id       665 non-null    int64    
 2   vol_requests     665 non-null    int64    
 3   event_time       665 non-null    int64    
 4   title            665 non-null    object    
 ..  ...              ...           ...  
 34  NTA              0 non-null    float64  
dtypes: float64(13), int64(8), object(14)  
memory usage: 182.0+ KB
```

- `object` : string/mixed types
- `int64` : integer
- `float64` : float
- `datetime64` : dates and times

Converting column types

```
print(df)
```

```
   A      B      C  
0 1  string  1.0  
1 2  string2 2.0  
2 3  string3 3.0
```

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 3 entries, 0 to 2  
Data columns (total 3 columns):  
 #   Column   Non-Null Count   Dtype     
 --  -----   -----    
 0   A         3 non-null     int64  
 1   B         3 non-null     object  
 2   C         3 non-null     object  
dtypes: int64(1), object(2)  
memory usage: 200.0+ bytes
```

Converting column types

```
print(df)
```

```
A      B      C  
0 1  string  1.0  
1 2  string2 2.0  
2 3  string3 3.0
```

```
df["C"] = df["C"].astype("float")  
print(df.dtypes)
```

```
A      int64  
B      object  
C      float64  
dtype: object
```

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Training and test sets

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman

Curriculum Manager, DataCamp

Why split?

1. Reduces *overfitting*
2. Evaluate performance on a holdout set

Splitting up your dataset

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
X_train y_train  
0      1.0    n  
1      4.0    n  
...  
5      5.0    n  
6      6.0    n
```

```
X_test y_test  
0      9.0    y  
1      1.0    n  
2      4.0    n
```

Stratified sampling

- Dataset of 100 samples: 80 **class 1** and 20 **class 2**
- Training set of 75 samples: 60 **class 1** and 15 **class 2**
- Test set of 25 samples: 20 **class 1** and 5 **class 2**

Stratified sampling

```
X_train,X_test,y_train,y_test = train_test_split(X, y, stratify=y, random_state=42)
```

```
y["labels"].value_counts()
```

```
class1    80  
class2    20  
Name: labels, dtype: int64
```

Stratified sampling

```
y_train["labels"].value_counts()
```

```
class1    60  
class2    15  
Name: labels, dtype: int64
```

```
y_test["labels"].value_counts()
```

```
class1    20  
class2     5  
Name: labels, dtype: int64
```

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON