

# Introduction to PySpark

INTRODUCTION TO PYSPARK



Benjamin Schmidt  
Data Engineer

# Meet your instructor

- Almost a Decade of Data Experience with PySpark
  - Used PySpark for Machine Learning, ETL tasks, and much more more
  - Enthusiastic teacher of new tools for all!
- 



# What is PySpark?

- Distributed data processing: Designed to handle large datasets across clusters
- Supports various data formats including CSV, Parquet, and JSON
- SQL integration allows querying of data using both Python and SQL syntax
- Optimized for speed at scale




# When would we use PySpark?

- Big data analytics
- Distributed data processing
- Real-time data streaming
- Machine learning on large datasets
- ETL and ELT pipelines
- Working with diverse data sources:
  1. CSV
  2. JSON
  3. Parquet
  4. Many Many More

# Spark cluster


## Master Node

- Manages the cluster, coordinates tasks, and schedules jobs



## Worker Nodes

- Execute the tasks assigned by the master
- Responsible for executing the actual computations and storing data in memory or disk



# SparkSession

- SparkSessions allow you to access your Spark cluster and are critical for using PySpark.

```
# Import SparkSession
from pyspark.sql import SparkSession

# Initialize a SparkSession
spark = SparkSession.builder.appName("MySparkApp").getOrCreate()
```

- `.builder()` sets up a session
- `getOrCreate()` creates or retrieves a session
- `.appName()` helps manage multiple sessions

# PySpark DataFrames

- Similar to other DataFrames but
- Optimized for PySpark

```
# Import and initialize a Spark session
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName("MySparkApp").getOrCreate()

# Create a DataFrame
census_df = spark.read.csv("census.csv",
                           ["gender", "age", "zipcode", "salary_range_usd", "marriage_status"])

# Show the DataFrame
census_df.show()
```

# **Let's practice!**

**INTRODUCTION TO PYSPARK**

# Introduction to PySpark DataFrames

INTRODUCTION TO PYSPARK



Benjamin Schmidt  
Data Engineer

# About DataFrames

- DataFrames: Tabular format (rows/columns)
- Supports SQL-like operations
- Comparable to a Pandas Dataframe or a SQL TABLE
- Structured Data



Name	Age	Country	Salary
Alice	35.0	USA	12000
Bob	30.0	USA	11000
Cathy	28.0	Canada	13000
David	32.0	Canada	14000
Eve	25.0	UK	10000
Fiona	30.0	UK	11500
Grace	22.0	Germany	13500
Hannah	28.0	Germany	14500
Ivan	35.0	France	15000
Jessica	25.0	France	12500
Karen	30.0	Spain	13500
Liam	28.0	Spain	14000
Mia	32.0	Australia	15500
Noah	26.0	Australia	14500
Olivia	30.0	Japan	16000
Parker	24.0	Japan	13000
Quinn	29.0	China	14000
Riley	31.0	China	15000
Sophia	27.0	India	12000
Taylor	33.0	India	13000
Ulysses	26.0	South Africa	11000
Vivian	30.0	South Africa	12000
Winston	28.0	Netherlands	14000
Xavier	32.0	Netherlands	15000
Yara	25.0	Sweden	13000
Zoe	30.0	Sweden	14000

# Creating DataFrames from filestores

```
# Create a DataFrame from CSV  
census_df = spark.read.csv('path/to/census.csv', header=True, inferSchema=True)
```

# Printing the DataFrame

```
# Show the first 5 rows of the DataFrame  
census_df.show()
```

	age	education.num	marital.status	occupation	income
0	90	9	Widowed		? <=50K
1	82	9	Widowed	Exec-managerial	<=50K
2	66	10	Widowed		? <=50K
3	54	4	Divorced	Machine-op-inspct	<=50K
4	41	10	Separated	Prof-specialty	<=50K

# Printing DataFrame Schema

```
# Show the schema  
census_df.printSchema()  
  
Output:  
root  
|-- age: integer (nullable = true)  
|-- education.num: integer (nullable = true)  
|-- marital.status: string (nullable = true)  
|-- occupation: string (nullable = true)  
|-- income: string (nullable = true)
```

# Basic analytics on PySpark DataFrames

```
# .count() will return the total row numbers in the DataFrame  
row_count = census_df.count()  
print(f'Number of rows: {row_count}')
```

```
# groupby() allows the use of sql-like aggregations  
census_df.groupBy('gender').agg({'salary_usd': 'avg'}).show()
```

Other aggregate functions are:

- `sum()`
- `min()`
- `max()`

# Key functions for PySpark analytics

- `.select()` : Selects specific columns from the DataFrame
- `.filter()` : Filters rows based on specific conditions
- `.groupBy()` : Groups rows based on one or more columns
- `.agg()` : Applies aggregate functions to grouped data

# Key Functions For Example

```
# Using filter and select, we can narrow down our DataFrame  
filtered_census_df = census_df.filter(df['age'] > 50).select('age', 'occupation')  
filtered_census_df.show()
```

Output

```
+-----+  
| age | occupation |  
+-----+  
| 90 | ? |  
| 82 | Exec-managerial |  
| 66 | ? |  
| 54 | Machine-op-inspct |  
+-----+
```

# **Let's practice!**

**INTRODUCTION TO PYSPARK**

# More on Spark DataFrames

INTRODUCTION TO PYSPARK



Benjamin Schmidt  
Data Engineer

# Creating DataFrames from various data sources

- CSV Files: Common for structured, delimited data
- JSON Files: Semi-structured, hierarchical data format
- Parquet Files: Optimized for storage and querying, often used in data engineering

- Example:

```
spark.read.csv("path/to/file.csv")
```

- Example:

```
spark.read.json("path/to/file.json")
```


- Example:

```
spark.read.parquet("path/to/file.parquet")
```

<sup>1</sup> [https://spark.apache.org/docs/latest/api/python/reference/pyspark.pandas/api/pyspark.pandas.read\\_csv](https://spark.apache.org/docs/latest/api/python/reference/pyspark.pandas/api/pyspark.pandas.read_csv)

# Schema inference and manual schema definition

- Spark can infer schemas from data with `inferSchema=True`
- Manually define schema for better control - useful for fixed data structures



# DataTypes in PySpark DataFrames

- `IntegerType` : Whole numbers
  - E.g., 1 , 3478 , -1890456
- `LongType`: Larger whole numbers
  - E.g., 8-byte signed numbers, 922334775806
- `FloatType` and `DoubleType`: Floating-point numbers for decimal values
  - E.g., 3.14159
- `StringType`: Used for text or string data
  - E.g., "This is an example of a string."
- ...

# DataTypes Syntax for PySpark DataFrames

```
# Import the necessary types as classes
from pyspark.sql.types import (StructType,
                                 StructField, IntegerType,
                                 StringType, ArrayType)

# Construct the schema
schema = StructType([
    StructField("id", IntegerType(), True),
    StructField("name", StringType(), True),
    StructField("scores", ArrayType(IntegerType()), True)
])

# Set the schema
df = spark.createDataFrame(data, schema=schema)
```

# DataFrame operations - selection and filtering

- Use `.select()` to choose specific columns
- Use `.filter()` or `.where()` to filter rows based on conditions
- Use `.sort()` to order by a collection of columns

```
# Select and show only the name and age columns  
df.select("name", "age").show()
```

```
# Filter on age > 30  
df.filter(df["age"] > 30).show()
```

```
# Use Where to filter match a specific value  
df.where(df["age"] == 30).show()
```

# Sorting and dropping missing values

- Order data using `.sort()` or `.orderBy()`
- Use `na.drop()` to remove rows with null values

```
# Sort using the age column  
df.sort("age", ascending=False).show()
```

```
# Drop missing values  
df.na.drop().show()
```

# Cheatsheet

- `spark.read_json()` : Load data from JSON
- `spark.read.schema()` : Define schemas explicitly
- `.na.drop()` : Drop rows with missing values
- `.select()` , `.filter()` , `.sort()` , `.orderBy()` : Basic data manipulation functions

# **Let's practice!**

**INTRODUCTION TO PYSPARK**